

Kong AI Gateway

Phased Team Org-Chart

Prepared for: Jesse Swensen, DXC

Date: 2026-05-12 (revised 2026-05-13)

Status: Draft — Internal Use Only

Scope: 130,000-employee enterprise AI traffic gateway

NEW FOUNDATIONAL DECISIONS — 2026-05-13

Two axes now govern the program. They are correlated but not coincident. Read both together to understand hiring, environment, and workload scope at any point in time.

Maturity tiers (Tier 1 / 2 / 3) describe the team shape — FTE count, coverage posture, governance model. The tier axis is the left column of every comparison in this document.

Implementation phases (Phase 1–4) describe gateway workload coverage — what AI traffic is in scope and what enforcement posture applies. Phases progress through Dev → Test → Prod under their own promotion gates, independently of which tier the team is at. See the Phase Plan section below for the full interaction.

ENVIRONMENT TIERS — DEV / TEST / PROD

Decision locked 2026-05-13. Every workload at every phase passes through these three environments in sequence under PR-gated promotion gates.

Dev	Test	Prod
PURPOSE Engineering iteration, plugin-chain experimentation, integration testing against mocked / sandbox model endpoints. Drift acceptable; never promotes. TRAFFIC Synthetic + sandbox provider endpoints. No real DXC employee traffic. OWNERS Gateway engineering (Role 1), DevOps (Role 2). Self-service for engineering. ON-CALL Best-effort, business hours. Engineer who last touched the relevant route.	PURPOSE Stakeholder UAT, security scanning (Garak, promptfoo, plugin-chain red-team), performance / load testing. Last stop before real users. TRAFFIC Mirrored / shadow traffic; controlled stakeholder cohorts; security and load harnesses. OWNERS Gateway engineering + AI-security (Role 5) + Compliance (Role 9) for UAT sign-off. ON-CALL Gateway practice owns during UAT windows. Working-hours escalation otherwise.	PURPOSE Live enterprise traffic for 130k DXC employees. Only env touching real provider keys with real cost attribution. TRAFFIC All real DXC AI traffic governed by the gateway. OWNERS Gateway engineering + on-call rotation (Role 12). Tier-1 NOC (Role 11) from Tier 2+. ON-CALL Only environment in the formal PagerDuty rotation. See on-call math below.

Standard promotion gate: Dev → Test → Prod. Every gateway-entity change (deck YAML) and every tenancy change (Terraform) passes through PR-gated `deck diff / terraform plan` review against each environment in sequence. No change reaches Prod that hasn't been verified in Test. Emergency hot-fixes get a 24h back-fill PR.

OPEN — ARCH TEAM CONFIRMS CONNECT ENV PATTERN (ROW 3)

PHASE PLAN — FOUR IMPLEMENTATION PHASES

Decision locked 2026-05-13. Phases describe workload coverage and enforcement posture. Each phase progresses independently through Dev → Test → Prod.

PHASE	POSTURE	WHAT LANDS	ENTRY GATE
Phase 1 — Discover	Observe-only. Gateway is a passive sensor.	Route all AI traffic to inventory it. File Log + OTel only — no enforcement plugins. Builds the risk inventory. Shadow AI surfaces here.	Arch sign-off + first Dev/Test/Prod control planes provisioned + first data plane deployed.
Phase 2 — Pareto risk reduction	Enforcing on the highest-risk 20% of traffic (80% of risk).	ai-prompt-guard, ai-sanitizer, ai-rate-limiting-advanced on real routes. Plugin-chain reviews per route mandatory. AI-security team becomes load-bearing.	Phase 1 risk inventory complete; AI-security policy authored; UAT in Test passes against shadow traffic.
Phase 3 — Agentic workloads	Enforcing on tool-use / MCP / multi-step agent traffic.	Routes for MCP servers, agent orchestrators, tool-using LLMs. Long-running session handling, multi-step trace correlation, per-tool authorization. Likely net-new tooling.	Phase 2 policy stable in Prod; agent workloads identified; MCP / agent expertise on team (see skill-gap note).

Phase 4 — Everything else	Full governance catch-all.	Remaining workloads under same plugin / observability / compliance regime. BU-liaison model is the scaling mechanism.	Phase 3 stable; federation model designed; embedded BU champions onboarded.
---------------------------	----------------------------	---	---

PHASE × TIER INTERACTION

PHASE	TYPICAL TIER	STAFFING IMPLICATION
Phase 1 — Discover	Tier 1 primary	Observe-only posture matches Tier 1 staffing reality. Observability lift (Role 3) is the dominant workload — pull Role 3 forward to 0.5 FTE in late Tier 1 rather than waiting for Tier 2 entry.
Phase 2 — Pareto risk	Tier 1 → Tier 2 transition	Enforcement on real traffic triggers dedicated AI-security headcount (Role 5 moves consumed → dedicated). Compliance partner becomes embedded. Phase 2 entry is the clearest Tier 1 → Tier 2 trigger.
Phase 3 — Agentic	Tier 2	MCP / agent / tool-use expertise is an explicit skill gap — none of Roles 1–13 carry it. Decision needed 3–4 months before Phase 3 entry: net-new hire, upskill Role 1 senior, or contract onramp specialist. See skill-gap callout below.
Phase 4 — Everything else	Tier 3	BU-liaison model (Role 13) is the scaling mechanism for onboarding the long tail. Central team focuses on platform health and federation governance.

PHASE × ENVIRONMENT-TIER INTERACTION

ENV	PHASE 1 — DISCOVER	PHASE 2 — PARETO RISK	PHASE 3 — AGENTIC	PHASE 4 — EVERYTHING
Dev	Synthetic / shadow traffic. Validate File Log + OTel pipeline config.	Plugin-chain experimentation against shadow traffic. Candidate plugin set tested.	MCP / agent route experimentation. New plugin compositions prototyped.	Catch-all route onboarding. BU champions operate Dev namespace for their BU.
Test	Stakeholder UAT on observe-only routes. Validate File Log + OTel pipelines.	Security and load harnesses against candidate plugin set. AI-security UAT sign-off.	Agent session load testing, per-tool authorization validation, trace correlation checks.	Full plugin / observability / compliance regime validated before Prod promotion.
Prod	Real DXC AI traffic, gateway-routed but unenforced. Risk inventory is built here.	Enforcement on the high-risk 20%. First live enforcing policies. FinOps attribution active.	Agent / MCP routes in production. New on-call runbooks for agent-session incidents.	All DXC workloads under full governance. Chargeback model covers entire portfolio.

Skill Gaps Surfaced by the Phase Plan For Nolan — Hiring Sequence

- ▲ **Phase 3 — MCP / agent / tool-use expertise.** None of Roles 1–13 explicitly carry this. Kong 3.x AI catalog does not yet have first-class MCP-aware audit or agent-session rate-limiting plugins — composition patterns are still being established. Decision needed 3–4 months before Phase 3 entry: (a) net-new senior hire with explicit MCP / agentic-workload experience — cleanest, slowest to source; (b) upskill an existing Role 1 senior with paid training time and an agent-platform mentor — cheapest, riskiest on timeline; (c) contract a Phase 3 onramp specialist for the first 6 months while option (a) or (b) is built — most flexible, highest cost.
- ▲ **Phase 1 — Observability-lift weighting.** Role 3 (Observability Engineer) is sized for Tier 2 steady state, but Phase 1 is observability-dominant by design. Tier 1 may need to over-index on observability earlier — pull Role 3 forward to 0.5 FTE in late Tier 1, or lean hard on the consumed DXC observability team during Phase 1. Nolan + Pax conversation on Role 3 sourcing recommended.
- ▲ **Env-tier coverage in Tier 1.** Three Konnect control planes (Dev / Test / Prod) is a 3× multiplier on configuration surface area at Tier 1. Current Tier 1 FTE (1.0 Role 1, 0.5–1.0 Role 2) may be tight. The Tier 1 → Tier 2 trigger of "P1 incident caused by single-point-of-failure" is more likely to fire earlier under the three-env model.

- Dedicated

FTE owned by the gateway practice
- Consumed / Consumed (earmarked)

FTE from adjacent DXC teams; earmarked = named contact + SLA-backed
- Governance

Architect / Tech Lead — technical direction & escalation
- Operations

NOC & on-call rotation — shared or rotation-based

TIER COMPARISON — GROWTH PATH AT A GLANCE

TIER 1 · ON-RAMP Pilot 0–6 months · First production routes	TIER 2 · TARGET STATE Production Enterprise Platform 6–18 months · Multi-region, multi-BU	TIER 3 · MATURE Federated Platform 18+ months · True enterprise product
--	--	--

Single region · 10–25 AI routes · 1–2 early-adopter BUs · Architecture decisions still landing · Business-hours coverage + best-effort after-hours

DEDICATED FTE 1.5–2.0	AI ROUTES 10–25
BUS SERVED 1–2	ON-CALL SLA Best-effort
CONSUMED PARTNERS (ACTIVE FROM DAY 1) AI Security · Compliance · Platform SRE · LLM Platform	
PHASE & ENV COVERAGE AT TIER 1 Phase 1 (Discover) is the primary workload. Phase 2 entry is the Tier 1 → Tier 2 trigger. Phases 3/4 not in scope. All three envs (Dev / Test / Prod) stand up at go-live. Sole gateway engineer owns ops across all three — 3× config surface is a declared risk.	

△ 1 gateway engineer = on-call every week. Unsustainable beyond 3–6 months. Plan Tier 2 within 12 months of go-live.

2–3 region data planes · 24/7 on-call SLA · Hundreds of AI routes · Dedicated AI-security team · FinOps dashboards · Embedded compliance partners

DEDICATED FTE 20–32	AI ROUTES 100s
BUS SERVED Multi-BU	ON-CALL SLA 24/7
CONSUMED / EARMARKED 3–5 FTE Platform/K8s SRE (earmarked) · Tier-1 NOC (shared service)	
PHASE & ENV COVERAGE AT TIER 2 Phase 2 (Pareto risk) at steady state across multi-BU traffic. Phase 3 (Agentic) entry happens during Tier 2 — MCP / agent skill gap must be resolved before Phase 3 entry. Flag for Nolan early in Tier 2 build-out. Dev / Test / Prod each scale to multi-region. Prod carries full FinOps + chargeback observability stack.	

Min 6 gateway engineers across 2–3 regions for credible 1-in-6 rotation. Below 4 globally = burnout territory.

Federation model · Central platform team + embedded BU liaisons per major BU · Published SLAs · Developer portal · Chargeback model · Red-team on quarterly cadence

TOTAL PRACTICE 30–50+	CENTRAL FTE ~20–30
BU LIAISONS 10–20+	ON-CALL SLA 1-in-8+
CONSUMED / EARMARKED 5–8 FTE Platform/K8s SRE (earmarked) · Tier-1 NOC (shared service)	
PHASE & ENV COVERAGE AT TIER 3 Phase 4 (Everything else) — BU-liaison model is the scaling mechanism for the long tail. Phases 1–3 all in steady state; new application discovery (Phase 1) runs continuously. BU champions hold merge rights to their BU's namespace in each env's control plane. Prod still under central on-call.	

BU liaisons resolve routine route issues before escalation — central rotation stays comfortable even past 20 BUs.

PER-TIER TEAM STRUCTURE

Tier 1 — Pilot
GOVERNANCE
Architect / Tech Lead Sets IaC patterns, owns Konnect tenancy decision (Option A vs B), escalation point0.2 FTE
DEDICATED CORE
Kong Gateway Engineer / SRE Sole gateway owner — deck YAML, Terraform (3 env control planes), AI plugin composition, on-call (best-effort). Must be staff/principal level.1.0 FTE
DevOps / GitOps Engineer CI pipelines, Helm chart mgmt, env promotion (Dev → Test → Prod gate). May be shared from platform pod early on.0.5–1.0 FTE
Product / Program Manager Defines AI routes needed, BU prioritization, internal SLA commitment0.2 FTE
CONSUMED (ACTIVE PARTNERS — DAY 1 MANDATORY)
AI-Aware Security Engineer Active in route onboarding, plugin reviews, threat modeling from day one — not a quarterly reviewer. Load-bearing at Phase 2 entry.Consumed
Compliance / Data-Privacy Officer Active from day one on any route touching regulated data. GDPR cannot be retrofitted.Consumed
Platform / Kubernetes SRE

Tier 2 — Production
GOVERNANCE
Architect / Tech Lead Full-time staff/principal. Technical roadmap, plugin-ecosystem, Konnect upgrades, federation model design1.0 FTE
DEDICATED CORE (GATEWAY ENGINEERING)
Kong Gateway Engineer / SRE Distributed across 2–3 regions for follow-the-sun. Non-negotiable floor for credible 24/7 SLA. Phase 3 MCP hire or upskill needed during this tier.6–10 FTE
DevOps / GitOps Engineer Platform automation, config pipelines (deck diff/Terraform for all envs), release management2–3 FTE
Observability Engineer OTel pipelines at billions-of-spans scale, FinOps dashboards, token-usage attribution per BU, anomaly detection2–3 FTE
DEDICATED — SECURITY, PLATFORM, COMPLIANCE & PM
AI-Aware Security Engineer Dedicated team — not a single engineer. Plugin-chain threat modeling, red-team cadence (Garak, promptfoo), guardrails-bypass dossier3–5 FTE
LLM Platform / Model Engineer Enterprise provider procurement, fallback orchestration, eval infrastructure for model-version upgrades. Per-env key lifecycle at hundreds-of-routes scale.1–2 FTE

Tier 3 — Federated
CENTRAL PLATFORM TEAM — GOVERNANCE
Architect / Tech Lead Technical roadmap, cross-BU standards, major platform decisions (Konnect major version, new plugin tier evaluation)1–2 FTE
CENTRAL PLATFORM TEAM — DEDICATED
Kong Gateway Engineers / SRE Smaller central team than T2 — BU liaisons absorb route-management volume. Focus: platform health, Konnect upgrades, plugin catalog6–8 FTE
DevOps / GitOps IaC repo automation, plugin test infrastructure, release management2–3 FTE
Observability FinOps dashboards, anomaly detection, capacity forecasting, SLO product2–3 FTE
AI-Security (+ Dedicated Red-Team) Policy authorship, OWASP LLM Top 10 / MITRE ATLAS cadence. 1–2 FTE explicit red-team on fixed quarterly cadence.3–5 + 1–2 FTE
LLM Platform / Model Provider procurement, fallback orchestration, eval infrastructure, plugin authorship for catalog gaps2–3 FTE

K8s namespace, cluster access, network policies, cert-manager for all three envs. Named contact required — not queue-shared.

Consumed (earmarked)

LLM Platform / Model Engineer

Provider API key lifecycle per env (sandbox / staging / prod keys). Joint key-rotation runbook.

Consumed

OPERATIONS

On-call (Prod only)

Gateway engineer carries on-call every week — declared risk, not a baseline. Dev/Test are best-effort working hours.

1 eng

Compliance / Data-Privacy Officer

GDPR, regional data-residency, sector-specific overlays (healthcare/finance/gov). Embedded — not fractional.

1-2 FTE

Product / Program Manager

Route roadmap, self-service portal, consumer onboarding, chargeback model, stakeholder narrative

2-3 FTE

CONSUMED / EARMARKED

Platform / Kubernetes SRE

3-5 FTE earmarked from DXC platform org. Named contacts per region. SLA-backed — P1/P2 in minutes, not days.

3-5 FTE earmarked

OPERATIONS

Tier-1 NOC

SLA-backed 24/7 first-line triage. Multi-region. Likely shared DXC NOC service.

Shared svc

Tier-2 On-call (PagerDuty rotation)

Drawn from gateway engineers (Prod only). Target: 1-in-6 rotation = ~8-9 on-call weeks/year per engineer.

6 eng min

Compliance

Central GDPR, EU AI Act, regional data-residency, sector-specific mapping. Per-BU overlays owned by BU compliance.

2-3 FTE

Product / Program Manager

Internal-product roadmap, SLA tiers, chargeback governance, BU champion certification program

2-3 FTE

FEDERATION LAYER — DEDICATED (EMBEDDED PER BU)

BU Gateway Champion / Liaison

1 per major BU, 2-3 per large BU. Owns route onboarding, BU-specific plugin config, first-line support. Has merge rights to BU namespace in each env's control plane. Trained & certified by central team.

10-20+ FTE

CONSUMED / EARMARKED

Platform / K8s SRE

Multi-region data-plane fleet. Dedicated named contacts per region.

5-8 FTE earmarked

OPERATIONS

Tier-1 NOC

Multi-region NOC with gateway-specific runbooks. BU liaisons resolve routine issues before escalation.

Shared svc

PHASE-TRANSITION TRIGGERS — "MOVE WHEN ANY TWO FIRE"

Tier 1

 →

Tier 2

 Target: 12 months after T1 go-live

Move when any TWO of the following are true:

› Route count crosses 25 production AI routes

› Consuming BU count reaches 3 or more

› P1 incident caused by single-point-of-failure (sole gateway engineer unavailable)

› On-call burnout documented in a retrospective

› A compliance requirement mandates 24/7 SLA

› Token-cost attribution becomes a real budget conversation (not a rounding error)

Tier 2

 →

Tier 3

 Target: 18+ months after T1 go-live

Move when any TWO of the following are true:

› BU count exceeds 15-20 distinct consuming business units (central onboarding bottleneck)

› On-call rotation requires >6 FTE to maintain health — signal that embedded BU capacity offloads route-management work better than more central hires

› Chargeback model has active BU negotiation on SLA customization — BUs want ownership, not service

› DXC M&A event adds a new organizational entity with a distinct compliance regime

ON-CALL MATH — GEOGRAPHIC DISTRIBUTION

Minimum to claim a 24/7 SLA honestly

6

GATEWAY ENGINEERS ACROSS 2-3 REGIONS

● Tier 1: 1 eng on-call every week. Declared risk — not a baseline.

● Tier 2: 1-in-6 rotation = ~8-9 on-call weeks/year per engineer. Comfortable.

● Tier 3: 1-in-8+ rotation. BU liaisons resolve routine issues before escalation.

Below 4 engineers globally: rotation degrades to 1-in-3 or worse — document as a staffing risk on any plan targeting fewer than 6 gateway engineers at Tier 2.

REGION	FTE (TIER 2 MIN)	PRIMARY COVERAGE WINDOW
Americas US-East or US-Central	2-3	08:00-20:00 ET primary On-call after-hours
EMEA UK or Continental EU	2	08:00-18:00 GMT/CET primary On-call after-hours
APAC India or APAC	1-2	08:00-18:00 IST/AEST primary On-call after-hours

